

a The document: one author, one quarter


one quarter, 2020Q2



...


Every ticker an author mentions in a quarter, across all posts and comments, is one document. Two tickers co-mention when a document holds both.

b An eligible pair: should have met, never did




Both tickers are frequent. Had authors picked tickers independently, at least 2 build-period documents would hold both. Observed: none. The pair is eligible. Every window has dozens to hundreds of such pairs.

c Rolling windows over 2019 to 2024



Eligible pairs are found in the build period. The statistic counts evaluation-period documents that hold any eligible pair. Build length was chosen by an outcome-blind census rule before any result.


d The statistic: observed count against a shuffle null



$z = (\text{observed} - \text{shuffle mean}) / \text{shuffle sd}$.
 z near 0: suppressed pairs meet as often as chance.
 z far below 0: kept apart (a wall). z above 0: pushed together. Each window's null uses its own documents.


e Three regimes, and the order $r/\text{wallstreetbets}$ passed through them

Chance-level mixing



through 2020Q3

Fusion



the two GameStop windows

Walls



2021Q2 onward, no reversion

Audiences of two suppressed tickers, drawn as the documents that mention each. Overlap is co-mention.