


a The document: one author, one quarter

one quarter, 2020Q2



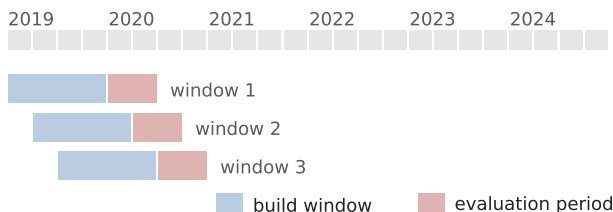
Every ticker an author mentions in a quarter, across all posts and comments, is one document. Two tickers co-mention when a document holds both.

b An eligible pair: should have met, never did



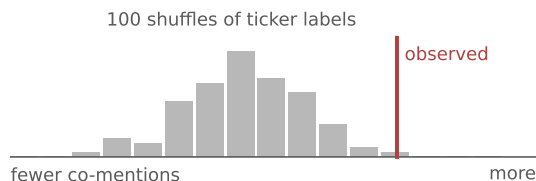
Both tickers are frequent. Independent choice predicts an expected count of at least 2 shared build documents. Observed: none. The pair is eligible; every window has dozens to hundreds of such pairs.

c Rolling windows over 2019 to 2024



Count document co-mentions separately for each eligible pair, then sum them. Build length follows the registered outcome-blind census rule.

d The statistic: observed count vs. label shuffles



$z = (\text{observed} - \text{shuffle mean}) / \text{shuffle SD}$.
Shuffled labels are deduplicated; margins can change.
Negative z : fewer joint mentions than under this null.
 z is an effect size, not a Gaussian tail probability.

e Three regimes under the original label null

